

Characterizing and subsetting Big Data workloads

Zhen Jia, Jianfeng Zhan , Lei Wang, Rui Han, Sally A. McKee, Qiang Yang,
Chunjie Luo, and Jingwei Li

1. Select representative workloads in big data fields
2. Identify a set of microarchitecture metrics that can directly or indirectly reflect program behavior
3. Standard statistical methods in characterizations

Identical Algorithms

- all selected applications have two implementations of the same algorithm
- offline analytics application: Hadoop & Spark
- interactive analytics application: Hive & Shark

Identical Data Sets

- For two applications, both data formats and data size are identical
- input data size for each workload is determined by the Spark
- For each workloads, data size ranging from several gigabytes to more than 100 gigabytes

Same Infrastructure

TABLE I. REPRESENTATIVE DATA ANALYSIS WORKLOADS

Category	Workload	Problem Size	Data Type	Software Stack
Offline Analytics	Sort	80 GB	unstructured sequence file	Hadoop & Spark
	WordCount	98 GB	unstructured text	
	Grep	98 GB	unstructured text	
	Naive Bayes	84 GB	semi-structured text	
	K-means	44 GB	unstructured text	
	PageRank	2^{24} vertices	unstructured graph	
Interactive Analytics	Projection	420 million records	structured table (e-commerce transaction data set)	Hive & Shark
	Filter	420 million records		
	Order By	420 million records		
	Cross product	100 million records		
	Union	420 million records		
	Difference	100 million records		
	Aggregation	420 million records		
	JoinQuery	100 million records		
	AggQuery	420 million records		
	SelectQuery	420 million records		

A broad set of metrics of different types that cover all major characters
---particularly focus on factors that may affect data movement or calculation

Nine Categories of 45 metrics

Eliminate correlated metrics before analysis

---Principle Component Analysis(PCA)

---only keep the top few PC's which have eigenvalues greater than or equal to one

Using Hierarchical clustering to analyze the similarity among workloads in two software stack

- Hierarchical clustering can quantitatively show the similarity among workloads via a dendrogram

Eight PCs


- retain 91.12% variance

All big data workloads(32)

Dendrogram

- drawing a U-shaped link between a non-singleton cluster and its children
- the length of the top of the U-link is the distance between its children
- the shorter the distance, the more similar the children
- linkage distance between two clusters is made by element pair, namely those two elements(one in each cluster) that are closest to each other

Methodology—Experiment A similarity analysis



Observation 1:

At first clustering iteration, 80% of clusters consist of workloads that are based on the same software stack.

Observation 2:


Two workloads implementing the same algorithm on the different software stacks do not get clustered together in the first clustering iteration- the only exception is Projection

Observation3:

After the first iteration, workloads based on the same software stack are easier to cluster together

Conclusion from Observation1&2&3:
software stacks have significant impacts on workload behaviors — even greater than that of the benchmark algorithms

Fig. 1. Similarity of Hadoop (H) and Spark (S) workloads.



Observation 4:
H-Union/H-Filter, S-Union/S-Filter, H-JoinQuery/H-CrossProduct, and S-JoinQuery/S-CrossProduct, are grouped, respectively. Workloads using the same software stack to implement similar algorithms have similar behaviors

Observation 5:
There are nine Hadoop workloads clustered within linkage distances of 2.72. Only three Spark-based workloads are clustered within a similar distance. Hadoop-based workloads have more similarities with each other than their Spark-based counterparts.

Conclusion from Observation 4&5:
Hadoop has a larger impact on microarchitectural behaviors than Spark

Fig. 1. Similarity of Hadoop (H) and Spark (S) workloads.

The first four PCs


---covering 71.45% variance

All big data workloads(32)

Scatter Plot

---Workloads appearing close in these figures may in fact be farther away when all PCs

--- workloads that appear far apart indeed exhibit significantly different behaviors




Spark-based workloads are spread widely along the x-axis (PC1) in Figure 2.

Hadoop-based workloads are grouped in the middle of the chart.

Conclusion: Spark-based workloads exhibit more diversity than their Hadoop-based counterparts with respect to PC1.

Fig. 2. Scatter plot of workloads using the first and second principle components



On the y-axis (PC2) we find that most of the Hadoop-based workloads are located towards the top, whereas most of Spark-based workloads are located in the middle or bottom parts of the chart.

Conclusion: PC2 is the main component that differentiates the Hadoop-based workloads from the Spark-based workloads

Fig. 2. Scatter plot of workloads using the first and second principle components

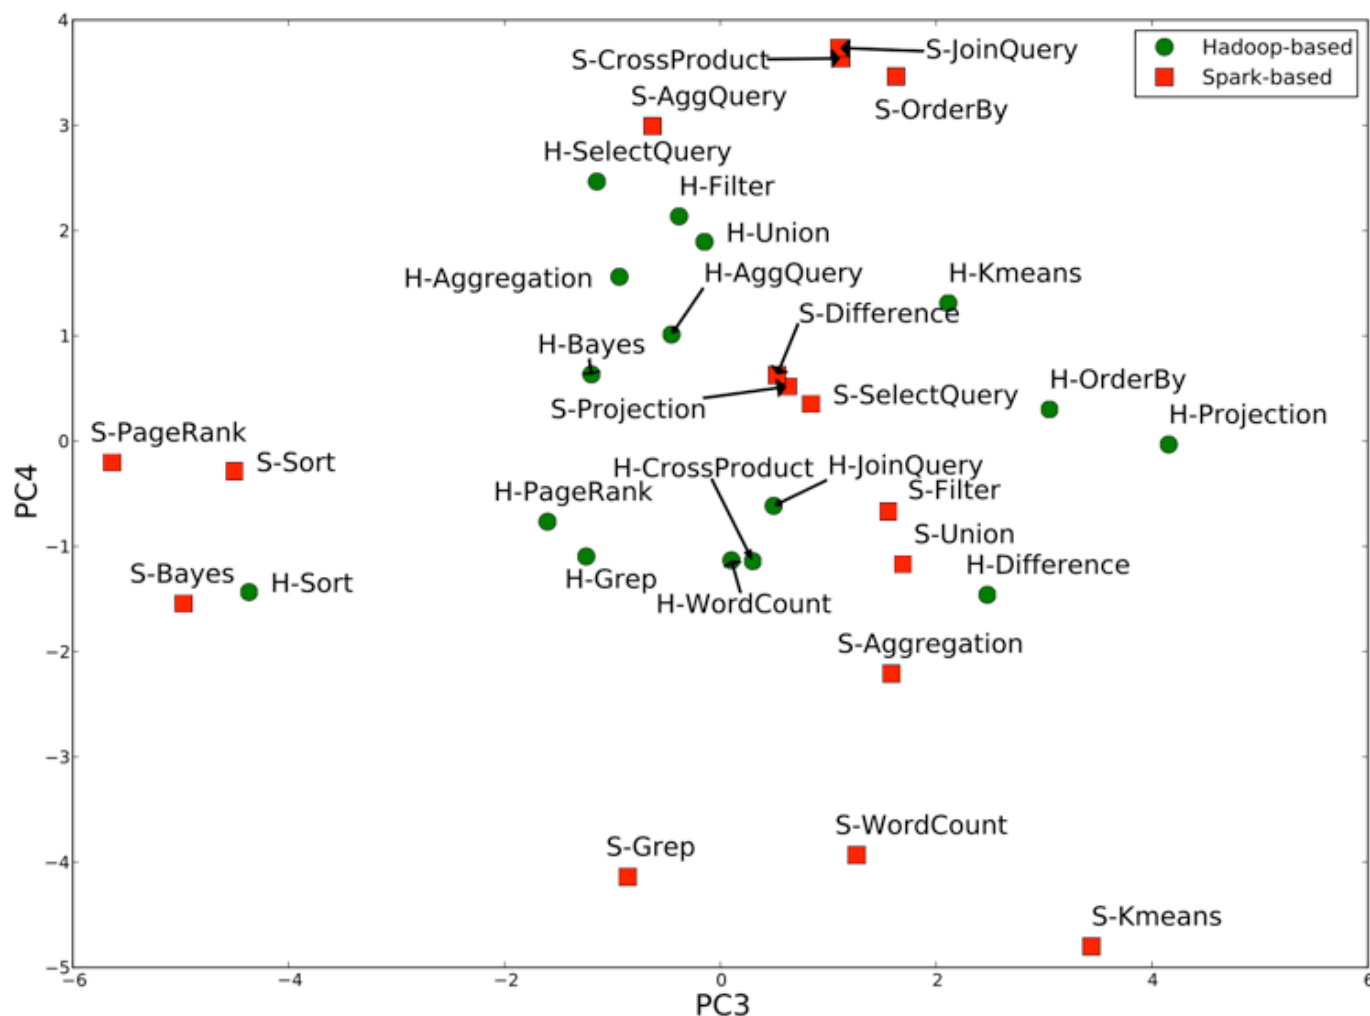


Fig. 3. Scatter plot of workloads using the third and fourth principle components

Spark-based workloads nearly cover all of the space for PC3 and PC4.

Hadoop-based workloads are again grouped in the center

Conclusion from Figure 2&3:

Verify the the result in experiment A


Spark-based workloads show more diversity than the Hadoop-based workloads with respect to microarchitectural behaviors because the Hadoop software stack minimizes the impact of the user application code

investigate which microarchitectural-level metrics dominate PC2 values

- Select 9 metrics negatively dominate PC2's value

- Select 5 metrics positively dominate PC2's value

- calculate the mean for each metric for each kind of big workload and present normalized Hadoop-based values using the Spark-based workloads as the baseline




Observation 6&7&8:

the large number of L3 cache and data TLB misses cause more backend stalls for the Spark-based workloads, and then high instruction cache misses cause more front-end stalls for the Hadoop-based workloads

Conclusion: Hadoop-based workloads have larger instruction footprints than their Spark-based counterparts. Spark-based workloads have larger data footprints than Hadoop-based workloads

Fig. 5. Metrics causing Hadoop and Spark to behave differently



Observation 9:
Spark-based workloads have more SNOOP HIT and SNOOP HITE responses than Hadoop-based workloads.

Conclusion:
Spark-based workloads have more data sharing among different cores.

Fig. 5. Metrics causing Hadoop and Spark to behave differently

Using eight principle components obtained in experiment A as metrics

K-means clustering to group the workloads

Choose a representative workload from each cluster

TABLE IV. THE RESULT OF K-MEANS CLUSTERING ALGORITHM

Cluster	Workloads	Number
1	H-PageRank, H-Grep, H-JoinQuery, H-Sort, H-CrossProduct, S-Bayes, S-Grep, S-Sort	8
2	H-AggQuery, S-Filter, S-Union, S-SelectQuery, S-OrderBy, H-Kmeans	6
3	S-Aggregation, S-WordCount, S-Kmeans, H-Wordcount, H-Bayes	5
4	H-OrderBy, H-Difference, S-Difference, S-PageRank	4
5	H-Aggregation, H-Filter, H-Union, H-SelectQuery	4
6	S-CrossProduct, S-JoinQuery, S-AggQuery	3
7	H-Projection, S-Projection	2

Two Method in K-means Clustering:

- choose the workload that is as close as possible to the center of the cluster it belongs to
- select an extreme workload situated at the “boundary” of each cluster

Using hierarchical clustering to select seven representative workloads, four of result workload s also included in the second approach in K-means while not in the the first approach

Conclusion: the second approach superior in selecting representative workloads

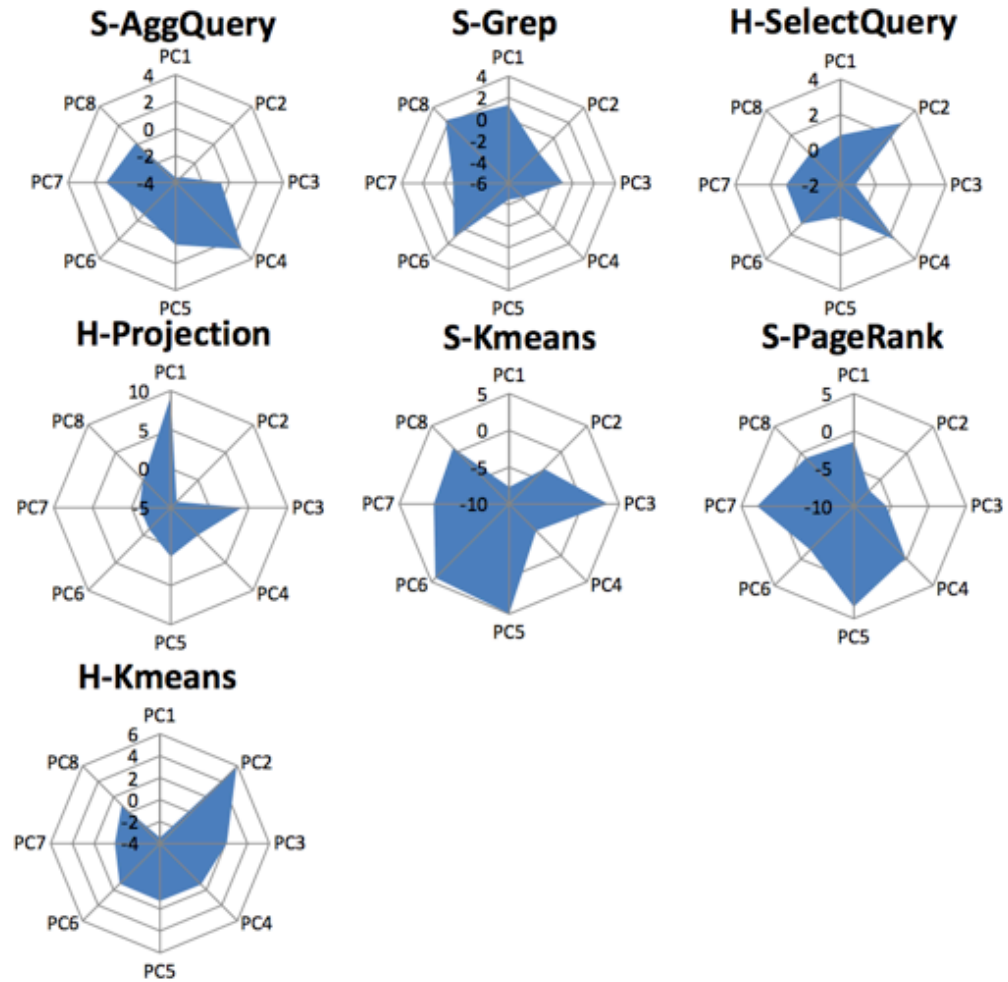


Fig. 6. Kiviat diagrams of the representative workloads

